

Yifan Zong

+1 (514) 585-8639
y22zong@uwaterloo.ca
github.com/zongyf02
zongyf02.github.io

Education

University of Waterloo, 3A Candidate for BSE, 4.0 GPA

2020–2025

Skills

Languages: Python, C++, C, SQL, R, Java, Swift, Lisp, JavaScript, Bash

Technologies: TensorFlow, JAX, Pytorch, Hugging Face, scikit-learn, Tidyverse, PySpark, gRPC

Experience

Cerebras Systems

ML Framework Developer

Sept. 2022–Present

Toronto, ON

TD Bank

Data Scientist

Jan.–Apr. 2022

Toronto, ON

- Implemented 4 variants of tri-training in **scikit-learn** and **TensorFlow** to research the application of semi-supervised learning in fraud detection.
- Benchmarked 18 tri-training models against supervised and self-trained baselines with positive results.
- Sped up training two-fold by applying data cleaning and feature selection on the cheque fraud dataset in **PySpark** and **scikit-learn**.

McMaster University

Data Scientist

Jan. 2021–Apr. 2022

Hamilton, ON

- Lead a 6-person team in the development of an **R** package for the cleaning and validation of the INORMUS study's datasets.
- Automated 3-hour data checks into 3-second functions with **Tidyverse** in **R**.
- Presented the **R** package to the McMaster Method Center and trained other researchers in its use.

Projects

[mlax](#)

Oct. 2022–Present

- Created a pure functional neural-network library built on top of **Google JAX** that is fully compatible with **JAX**'s native transformations.
- Implemented linear, convolution layers, and more, using the `jax.lax` package.
- Trained an MLP on MNIST that matched the accuracy of a **Pytorch** reference implementation using a quarter of the training time.

[Chess AI](#)

Jul.–Aug. 2022

- Created a NegaMax AI using PeSTO's evaluation function in **C++**.
- Implemented move-ordering, PV-moves, killer moves, and history heuristics to improve the effectiveness of alpha-beta pruning.
- Achieved a practical search depth of 6 and an ELO of ~1000.

[NameThatFish](#)

May–Jun. 2022

- Created custom image augmentation layers in **Tensorflow** and fine-tuned an EfficientNetV2 model to classify scraped images of 20 Ontarian fish species.
- Achieved 90+% average accuracy on test data.
- Deployed the model on **Hugging Face Spaces** using the **Gradio** API.